

Lecture 3. tidyverse-ggplot2

Lazaro Alonso

6/2/2020

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
summary(cars)
```

```
##      speed      dist
##  Min.   : 4.0    Min.   :  2.00
## 1st Qu.:12.0    1st Qu.: 26.00
##  Median:15.0    Median : 36.00
##   Mean  :15.4    Mean   : 42.98
## 3rd Qu.:19.0    3rd Qu.: 56.00
##   Max.  :25.0    Max.   :120.00
```

Including Plots

You can also embed plots, for example:



Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.

Data distributions

Let's see one more time the mpg data set.

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.0 --
## v ggplot2 3.3.0      v purrr  0.3.4
## v tibble  3.0.1      v dplyr  0.8.5
## v tidyr   1.1.0      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.5.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

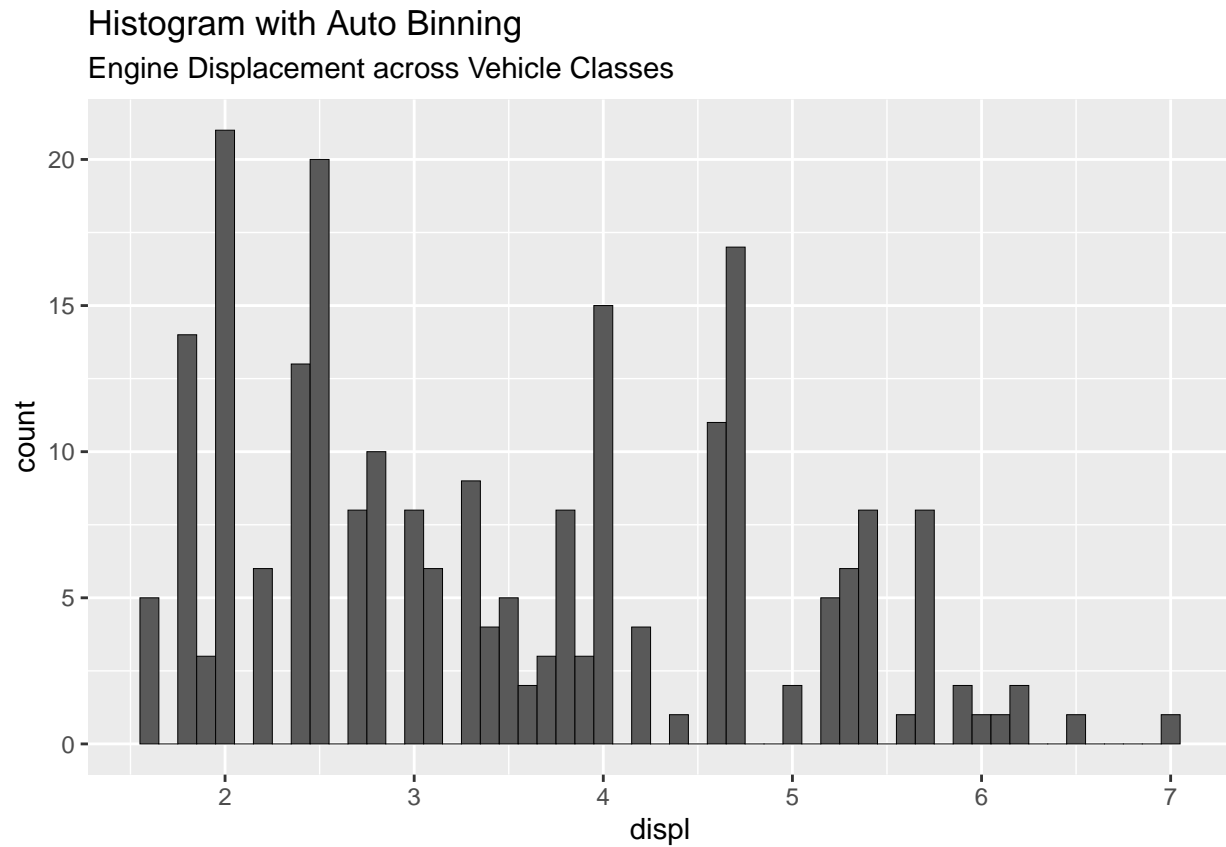
glimpse(mpg)

## Rows: 234
## Columns: 11
## $ manufacturer <chr> "audi", "audi", "audi", "audi", "audi", "audi", "audi"...
## $ model        <chr> "a4", "a4", "a4", "a4", "a4", "a4", "a4", "a4 quattro"...
## $ displ        <dbl> 1.8, 1.8, 2.0, 2.0, 2.8, 2.8, 3.1, 1.8, 1.8, 2.0, 2.0,...
## $ year         <int> 1999, 1999, 2008, 2008, 1999, 1999, 2008, 1999, 1999, ...
## $ cyl          <int> 4, 4, 4, 6, 6, 6, 4, 4, 4, 6, 6, 6, 6, 6, 8, ...
## $ trans        <chr> "auto(l5)", "manual(m5)", "manual(m6)", "auto(av)", "a...
## $ drv          <chr> "f", "f", "f", "f", "f", "f", "f", "4", "4", "4", "4",...
## $ cty          <int> 18, 21, 20, 21, 16, 18, 18, 18, 16, 20, 19, 15, 17, 17...
## $ hwy          <int> 29, 29, 31, 30, 26, 26, 27, 26, 25, 28, 27, 25, 25, 25...
## $ fl           <chr> "p", "p", "p", "p", "p", "p", "p", "p", "p", "p", "p",...
## $ class        <chr> "compact", "compact", "compact", "compact", "compact",...

help(mpg)

g <- ggplot(mpg, aes(displ)) + scale_fill_brewer(palette = "Spectral")

g + geom_histogram(binwidth = .1,
                   col="black",
                   size=.1) + # change binwidth
  labs(title="Histogram with Auto Binning",
       subtitle="Engine Displacement across Vehicle Classes")
```




Let's do it per class.

```
library(tidyverse)
g <- ggplot(mpg, aes(displ)) + scale_fill_brewer(palette = "Spectral")

g + geom_histogram(aes(fill=class),
  binwidth = .1,
  col="black",
  size=.1) + # change binwidth
labs(title="Histogram with Auto Binning",
  subtitle="Engine Displacement across Vehicle Classes")
```


Histogram with Auto Binning
Engine Displacement across Vehicle Classes



```
g + geom_histogram(aes(fill=class),
  bins=10,
  col="black",
  size=.1) + # change number of bins
labs(title="Histogram with Fixed Bins",
  subtitle="Engine Displacement across Vehicle Classes")
```

Histogram with Fixed Bins

Engine Displacement across Vehicle Classes




Histogram on a categorical variable

```
g <- ggplot(mpg, aes(manufacturer))
g + geom_bar(aes(fill=class), width = 0.5) +
  theme(axis.text.x = element_text(angle=65, vjust=0.6)) +
  labs(title="Histogram on Categorical Variable",
        subtitle="Manufacturer across Vehicle Classes")
```

Histogram on Categorical Variable

Manufacturer across Vehicle Classes




Density plot

```
g <- ggplot(mpg, aes(cty))
g + geom_density(aes(fill=factor(cyl)), alpha=0.5) +
  labs(title="Density plot",
        subtitle="City Mileage Grouped by Number of cylinders",
        caption="Source: mpg",
        x="City Mileage",
        fill="# Cylinders")
```

Density plot

City Mileage Grouped by Number of cylinders




Source: mpg

Box plots

The dark line inside the box represents the median. The top of box is 75%ile and bottom of box is 25%ile. The end points of the lines (aka whiskers) is at a distance of $1.5 \times \text{IQR}$, where IQR or Inter Quartile Range is the distance between 25th and 75th percentiles. The points outside the whiskers are marked as dots and are normally considered as extreme points.

```
g <- ggplot(mpg, aes(class, cty))
g + geom_boxplot(varwidth=T, fill="dodgerblue", alpha = 0.5) +
  labs(title="Box plot",
        subtitle="City Mileage grouped by Class of vehicle",
        caption="Source: mpg",
        x="Class of Vehicle",
        y="City Mileage")
```




Source: mpg

Setting `varwidth=T` adjusts the width of the boxes to be proportional to the number of observation it contains.

```
g <- ggplot(mpg, aes(class, cty))
g + geom_boxplot(aes(fill=factor(cyl))) + scale_fill_brewer(palette = "Spectral") +
  theme(axis.text.x = element_text(angle=65, vjust=0.6)) +
  labs(title="Box plot",
        subtitle="City Mileage grouped by Class of vehicle",
        caption="Source: mpg",
        x="Class of Vehicle",
        y="City Mileage")
```


Box plot

City Mileage grouped by Class of vehicle




Source: mpg

Violin plot

```
g <- ggplot(mpg, aes(class, cty))
g + geom_violin() +
  labs(title="Violin plot",
        subtitle="City Mileage vs Class of vehicle",
        caption="Source: mpg",
        x="Class of Vehicle",
        y="City Mileage")
```

Violin plot

City Mileage vs Class of vehicle



Source: mpg

```
g <- ggplot(mpg, aes(class, cty))
p <- g + geom_violin() +
  labs(title="Violin plot",
        subtitle="City Mileage vs Class of vehicle",
        caption="Source: mpg",
        x="Class of Vehicle",
        y="City Mileage")
p
```

Violin plot


City Mileage vs Class of vehicle



Source: mpg

```
g + geom_violin(TRIM=FALSE) + coord_flip()
```

```
## Warning: Ignoring unknown parameters: TRIM
```



Add summary statistics on a violin plot

```
p + stat_summary(fun.y=median, geom="point", size=2, color="red")
```

Warning: `fun.y` is deprecated. Use `fun` instead.

Violin plot

City Mileage vs Class of vehicle




Source: mpg

```
## Add median and quartile  
p + geom_boxplot(width=0.1)
```

Violin plot

City Mileage vs Class of vehicle




Source: mpg

```
g <- ggplot(mpg, aes(class, cty, color = class))
pc <- g + geom_violin() +
  labs(title="Violin plot",
        subtitle="City Mileage vs Class of vehicle",
        caption="Source: mpg",
        x="Class of Vehicle",
        y="City Mileage")
pc
```

Violin plot

City Mileage vs Class of vehicle




Source: mpg

```
pc + scale_color_brewer(palette="Dark2")
```

Violin plot

City Mileage vs Class of vehicle



Source: mpg

More than one plot

```
# install.packages("gridExtra")
library("gridExtra")
```

```
##
```

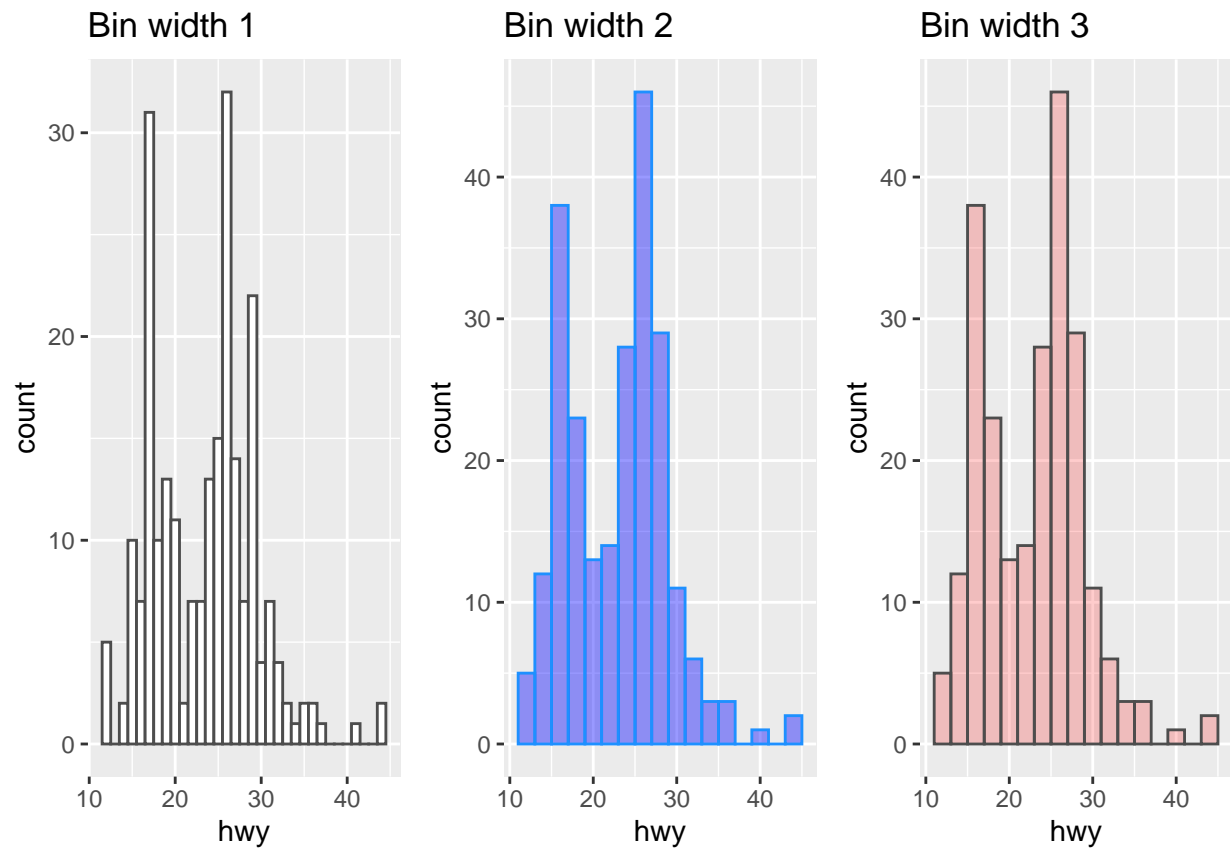
```
## Attaching package: 'gridExtra'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
## combine
```

```
p_data <- ggplot(mpg, aes(x = hwy))
p1 <- p_data + geom_histogram(binwidth = 1, color = "grey30", fill = "white") +
  labs(title="Bin width 1")
p2 <- p_data + geom_histogram(binwidth = 2, color = "dodgerblue",
                              fill = "blue", alpha = 0.4) +
  labs(title="Bin width 2")
p3 <- p_data + geom_histogram(binwidth = 2, color = "grey30",
                              fill = "red", alpha = 0.2) +
  labs(title="Bin width 3")
grid.arrange(p1, p2, p3, ncol = 3)
```





overlap density and histogram

```
ggplot(mpg, aes(x = hwy)) + geom_histogram(aes(y=..density..), color = "grey30", fill = "white") + geom_density(aes(y=..density..), color = "red", fill = "white")
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

Bin width 1




Add Value Markers

```
ggplot(mpg, aes(x = hwy)) + geom_histogram(aes(y=..density..), color = "grey30", fill = "white") + geom_
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Bin width 1




To add lines for grouped data we need to do a little computation prior to plotting.

```
compare_mean <- mpg %>%
  group_by(cyl) %>%
  summarise(Mean = mean(cty))

g <- ggplot(mpg, aes(cty))
g + geom_density(aes(fill=factor(cyl)), alpha=0.35) +
  labs(title="Density plot",
        subtitle="City Mileage Grouped by Number of cylinders",
        caption="Source: mpg",
        x="City Mileage",
        fill="# Cylinders") +
  geom_vline(data = compare_mean, aes(xintercept = Mean, color = factor(cyl)),
            linetype = "dashed", size = 0.75)
```

Density plot

City Mileage Grouped by Number of cylinders



Source: mpg

Error bars


```
# Create new dataframe mpg_means_se
mpg_means_se <- mpg %>%
  group_by(manufacturer) %>% # Group the data by manufacturer
  summarize(mean_cty=mean(cty), # Create variable with mean of cty per group
            sd_cty=sd(cty), # Create variable with sd of cty per group
            N_cty=n(), # Create new variable N of cty per group
            se=sd_cty/sqrt(N_cty), # Create variable with se of cty per group
            upper_limit=mean_cty+se, # Upper limit
            lower_limit=mean_cty-se # Lower limit
  )
mpg_means_se
```

A tibble: 15 x 7

	manufacturer	mean_cty	sd_cty	N_cty	se	upper_limit	lower_limit
	<chr>	<dbl>	<dbl>	<int>	<dbl>	<dbl>	<dbl>
## 1	audi	17.6	1.97	18	0.465	18.1	17.1
## 2	chevrolet	15	2.92	19	0.671	15.7	14.3
## 3	dodge	13.1	2.49	37	0.409	13.5	12.7
## 4	ford	14	1.91	25	0.383	14.4	13.6
## 5	honda	24.4	1.94	9	0.648	25.1	23.8
## 6	hyundai	18.6	1.50	14	0.401	19.0	18.2
## 7	jeep	13.5	2.51	8	0.886	14.4	12.6
## 8	land rover	11.5	0.577	4	0.289	11.8	11.2
## 9	lincoln	11.3	0.577	3	0.333	11.7	11
## 10	mercury	13.2	0.5	4	0.25	13.5	13
## 11	nissan	18.1	3.43	13	0.950	19.0	17.1


```
## 12 pontiac      17    1      5 0.447      17.4      16.6
## 13 subaru      19.3 0.914    14 0.244      19.5      19.0
## 14 toyota      18.5 4.05     34 0.694      19.2      17.8
## 15 volkswagen  20.9 4.56     27 0.877      21.8      20.0
```

```
ggplot(mpg_means_se, aes(x=manufacturer, y=mean_cty, fill=manufacturer)) +
  geom_errorbar(aes(ymin=lower_limit, ymax=upper_limit), width = 0.3) +
  geom_bar(stat="identity", alpha= 0.5) +
  theme(axis.text.x = element_text(angle=65, vjust=0.6), legend.position = "none") +
  geom_line(aes(group = 1), size = 0.2, color = "red")
```



Reordering...

```
ggplot(mpg_means_se, aes(x=reorder(manufacturer, mean_cty), y=mean_cty, fill=manufacturer)) +
  geom_errorbar(aes(ymin=lower_limit, ymax=upper_limit), width = 0.3) +
  geom_bar(stat="identity", alpha= 0.5) +
  theme(axis.text.x = element_text(angle=65, vjust=0.6), legend.position = "none") +
  geom_line(aes(group = 1), size = 0.2, color = "red")
```



Now, use all these examples to explore the nycflights data set.

EXAMPLE:

```
library(nycflights13)
head(flights)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Data sets

```
ls("package:nycflights13")
```

```
## [1] "airlines" "airports" "flights"  "planes"   "weather"
```

```
airlines_data <- airlines
airports_data <- airports
flights_data <- flights
```

```
planes_data <- planes
weather_data <- weather
# Variables in flights dataset
?flights
```

```
head(weather)
```

```
## # A tibble: 6 x 15
##   origin year month   day hour temp dewp humid wind_dir wind_speed wind_gust
##   <chr>   <int> <int> <int> <int> <dbl> <dbl> <dbl>   <dbl>    <dbl>    <dbl>
## 1 EWR     2013     1     1     1  39.0  26.1  59.4     270     10.4      NA
## 2 EWR     2013     1     1     2  39.0  27.0  61.6     250      8.06     NA
## 3 EWR     2013     1     1     3  39.0  28.0  64.4     240     11.5      NA
## 4 EWR     2013     1     1     4  39.9  28.0  62.2     250     12.7      NA
## 5 EWR     2013     1     1     5  39.0  28.0  64.4     260     12.7      NA
## 6 EWR     2013     1     1     6  37.9  28.0  67.2     240     11.5      NA
## # ... with 4 more variables: precip <dbl>, pressure <dbl>, visib <dbl>,
## #   time_hour <dtm>
```

```
ggplot(weather, aes(x = temp)) +
  geom_histogram(binwidth = 3, color = "grey", fill = "red", alpha= 0.5) +
  labs(x = "Temperature (degrees F)", y = "Count",
       title = "New York City Airport Temperatures 2013")
```

```
## Warning: Removed 1 rows containing non-finite values (stat_bin).
```

